

Meta-Llama-3-8B-Instruct | input=4096, output=256 | Throughput vs Batch Size

